

Introduction

This is the first mandatory assignment for the course SIGB F2013. In this assignment we'll implement a simple gaze tracker. This will be done in the programming language python with help from the opencv and numpy libraries.

This report is an attempt to document what has been done to make this gaze tracker.

The basic structure for each section will be a short introduction of the goal of the section, followed by the theory behind our approach ending with a description of our actual implementation with visual aids used for documentation. Additionally we will accompany the report with captured videos demonstrating the usage of the eye tracker




Figure 1: Eye located with hough

Pupil detection

Overall rationale and theory

The eye consists of several distinct features such as pupil, iris, limbus and sclera. Furthermore several features are present in images as result of the light conditions when a picture is taken. The position of these can aid in correctly determining the gaze. Detecting these features is therefore a good starting for a gaze tracker implementation, but it poses some challenges in correctly identifying each component, and filtering away noise.

Of these features one of the easiest to recognize is the pupil, as it is the darkest part of the eye. Furthermore it's good starting point, as it is surrounded by the other eye components. The challenges/noise in locating the pupil is mostly due to glints/reflections of light, but these also aids in locating the pupil because we know that the pupil will reflect light in a certain way.

As mentioned, the pupil is the darkest part of the eye. Furthermore a pupil will reflect two “glints” of light. A good way to find the pupil is to find an intensity value which can separate it from the background. This is called thresholding, and will be explained in the following sections. Thresholding can be applied to both find pupil candidates as well as glints used for more robustly picking the best pupil candidate. When thresholding has been performed we will analyse certain features of the BLOBs in the image to see which are good candidates for pupil and glints. To aid in selecting a proper threshold we will use a form of pixel classification to automatically set a threshold. This method will be described last in this section.

Thresholding

Thresholding is a form of point processing used for separating areas of an image based on intensity in these areas. The desired end result of this method is a binary image with the foreground (object) in one color and the background (everything else) in a different color. This is usually black and white respectively. This will effectively separate the foreground from the background for us, leaving us with only the contours. In that sense we lose information and granularity in the image, but since we’re only interested in position we haven’t lost anything important.

Theory

As with all point processing methods the basic theory behind thresholding is applying a calculation to every pixel in the image, effectively changing the value of that pixel. We know that the end result is a binary image, following this, each pixel will be transformed into one of two values. We’ve so far operated on byte images so the max value is 255 and the min value is of course 0. For clarity we will assign each pixel either the min or the max value. Thus thresholding can be expressed as the following with T being the assigned Threshold value and f being the function applied to each pixel:

$$\text{if } f(x, y) \leq T \text{ then } g(x, y) = 0 \quad \text{and} \quad \text{if } f(x, y) > T \text{ then } g(x, y) = 255$$

Performing these operations will leave us with an image where only certain BLOBs are visible. We can then analyse the properties of these BLOBs to find which one is most likely to be a pupil, and which are most likely to be glints. Namely we will look at the area and the circularity of blobs. Area is simply a count of all pixels in the blob.

Finding the circularity is a bit more complex. Used here is “Heywoods circularity factor” which is derived from the perimeter and area of the BLOB. Perimeter is the count of pixels on the rim of the contour. A “cheaper” approximation can be found by taking the perimeter of the bounding box of a BLOB. The

bounding box can be found simply by finding the lowest and highest values of x and y respectively within a BLOB.

Once we have the perimeter and the area heywoods method can be applied. It's defined as follows:

$$Circularity = \frac{perimeter}{2*\sqrt{\pi*area}}$$

The area and circularity will be used to identify the best candidates

Our implementation

The usefulness of thresholding relies on having a total binary image with the foreground easily distinguishable from the background. The foreground is the pupil. The pupil is a dark circle-like object surrounded by a lighter circle (the iris). So we're looking for a threshold value that is lower than the surrounding iris. In an ideal world the pupil in it's entirety will have only one intensity throughout, and the ideal threshold value will be that. However, the world is rarely so black and white (literally) and this also isn't the case here. Although the pupil is the darkest spot it's not completely dark. So a threshold value of 1 would be far to low in most circumstances. It also requires very specialized lighting conditions to achive a pupil with the same intensity throughout, so there is no "silver bullet" threshold value which will perfectly cover the entire pupil.

What we're looking for is a value that is close enough to every part of the pupil and far enough away from every part of the iris. The way of finding this perfect value is mostly trial and error in this stage. Luckily we had a slider to play around with when searching for this value. For the thresholding itself, a built in cv2 method was used.

```
val,binI =cv2.threshold(gray, thr, 255, cv2.THRESH_BINARY_INV)
```

Where:

gray is our (grayscale) image

thr is our selected threshold value

255 is the maximum value

the last argument is a constant indicating that output is a binary image

We quickly saw that a good value for most of the sequences floats around 100. An example of this can be seen in figure 2:

This value is not robust in all cases, and too high or too low values will yield to few or too many results respectively, seen in figures 3 and 4

Another aspect of this is that further analysis is needed on the contours. So if a distorted figure is all we have it will be difficult or impossible to correctly




Figure 2: Good threshold




Figure 3: Low threshold




Figure 4: High threshold

determine the nature of the contour. An example of this false classification is seen in figure 5




Figure 5: Contours found with threshold

In the happy path scenario analysis of the blobs will be performed as described in the theory section above. First the area of the BLOB is found:

This will be found using a cv2 methods. Presumably it's naively implemented and has linear complexity, but this is speculation as we don't have the implementation. Usage is as follows:

```
a = cv2.contourArea(con)
```

Where con is a contour

Next the perimeter is found. Again opencv has an implementation for this, given a closed contour:

```
p = cv2.arcLength(con, True)
```

With these variables in hand the results can be filtered in a simple imperative manner:

```
if(a==0 or a<minArea or a>maxArea):
    continue
p = cv2.arcLength(con, True)
m = p/(2.0*math.sqrt(math.pi * a))
if (m<1.7):
```

```

if(len(con)>=5):
    ellips = cv2.fitEllipse(con)
    matches.append(ellips)

```

The min and max area as well as the circularity value of 1.7 are found through trial and error. The area parameters can be adjusted using sliders as needed for each sequence. The constraints of the length of the con is because we need 5 parameters for an ellipse

Last step is to redo this with the glints.




Figure 6: Glints with threshold

When we have the two glints of the pupil we can filter the data further based on these newly found features. Our algorithm for this is very straight forward and imperative:

```

for candA in glints:
    for candB in glints:
        #only accepting points with a certain distance to each other.
        if (Distance(candA[0],candB[0])> sliderVals['glintMinDist'] and Distance(candA[0],candB[0])< sliderVals['glintMaxDist']):
            glintList.append(candA)

#run through the remaining glints, keeping those that are close to the pupil candidates.
for glintCand in glintList:
    for pupCand in pupils:
        if(Distance(glintCand[0],pupCand[0])>sliderVals['glint&pubMINDist'] and Distance(glintCand[0],pupCand[0])< sliderVals['glint&pubMAXDist']):
            glintList1.append(glintCand)

#run through the pupil candidates keeping those that are close to the final glints list

```

```

for candP in pupils:
    for glintCand in glintList1:
        if(Distance(candP[0],glintCand[0])>sliderVals['glint&pubMINDist'] and Distance(candP[0],glintCand[0])>sliderVals['pupilMINDist']):
            pupilList.append(candP)

#sort out the pupils too far away from the found glints.
return (set(glintList1),set(pupilList))

```

One critique of this approach could be that the pupil and glints depend on each other “both ways”. That is, we first filter the glint candidates based on the pupil candidates and then the reverse is done. However, the results seems to be fairly precise

Pixel classification

In this section it will be demonstrated how a form of machine learning is applied to perform supervised classification in order to semi-automatically set a correct thresholding value. Correct is defined as a value which will allow us to perform the steps described in the previous section

Clustering is the practice of grouping a set of elements into several smaller groups of elements with similar features. It has a wide range of applications within datamining and different types of analysis. Obviously what will be demonstrated here is it's use within image analysis. Clustering isn't a specific algorithm. Rather it's the task we wish to perform in order to achieve our goal. In this assignment we've used the k-means algorithm for this

K-means is a clustering algorithm which can group a number of observations/data points into k number of clusters based on their nearest mean value.

The properties of the k-means algorithm (k groups based on mean intensity) combined with an existing knowledge of the properties of the eye(the pupil is the darkest part) makes it possible to use k-means for setting a threshold value automatically

Theory

The basic idea behind k-means clustering is to iteratively run through a dataset assigning points in their correct cluster based on previously selected values. Initially k points are selected and denoted as a center for it's cluster, c_1, \dots, c_k . These points can be selected on random or based on some guessed distribution. On each run through the dataset every point is examined. For each point the closest c is found, and the point is marked as to belong to this cluster. Once all points have been examined and placed in a cluster, the mean value of each

cluster is calculated as $c_i val$. Compare the mean value for the cluster with the previously recorded value of $c_i val$. If it has changed, another run through is performed. This continues until a desired level of precision is achieved or amount of runs have taken place

When we have done this we have k different threshold values to choose from, given our knowledge of the pupil, we will choose the one with the lowest mean value ($c_i val$).

Our implementation

There are a couple of possible caveats for this approach. Firstly there is an element of uncertainty in exactly how the clusters will be distributed. We need therefore to have a high enough k value to be sure to get the right cluster. There is also some uncertainty about whether the pupil always belongs to the darkest cluster. If for instance our k -value is too high, and there exists a darker region in the image (dark spot on the skin for instance) then the value of this cluster will be chosen as a threshold value, and we might miss the pupil

Through trials it was found that 8 is a good value for k in the sense that it often proved to segment the picture enough to allow the intensity value of the pupil to exist in one of the clusters. The other parameter is a constant that we apply so that the distance between the points has less importance by dividing each point with this constant. For this 15 was chosen as a good value.

The first step is performance optimization by resizing the image. This saves a lot of computations:

```
smallI = cv2.resize(gray, reSize)
```

where `reSize` is a 1x2 vector (30x30 was used)

When this is done we operate on column vectors

```
M,N = smallI.shape
X,Y = np.meshgrid(range(M),range(N))

z = smallI.flatten()
x = X.flatten()
y = Y.flatten()
O = len(x)
```

From this we can create a “feature” matrix with values corresponding to the intensity value of each pixel. Then the k -means algorithm is applied to this feature-matrix to produce the centroids of each cluster. The specific k -means implementation comes from `scipy`. The same library also provides a vector

quantization which gives us the label for each feature. Combining this we're left with a label image with intensities. From these we pick the darkest and take that value as to mean the intensity of the pupil. In reality this could probably be analysed further to more robustly identify which cluster contains the pupil. For instance we could analyse the original image and verify that there exists an area with this intensity which looks sufficiently like a pupil. If not, the second darkest area could be explored. This will however further increase the computational complexity.




Figure 7: K-Means

Usefulness of the k-Means algorithm

K-Means seems very robust and definitely useful in applications like this. A huge downside to this however is that it's computationally expensive. Our implementation is at least. To overcome this we've used a downsized image. As can be seen in figure ???. This also makes it very hard to visualize the method, as 30x30 pictures doesn't look to great. If the downsizing isn't done, then the sequence can't really run. One could argue that the lighting conditions of a picture rarely changes on each frame. So doing k-means every time the frame is updated might be overkill, and some calculations could be saved.